

# Memory Management

Sirak Kaewjamnong

517-312 Operating Systems

# Memory hierarchy

Computers have:

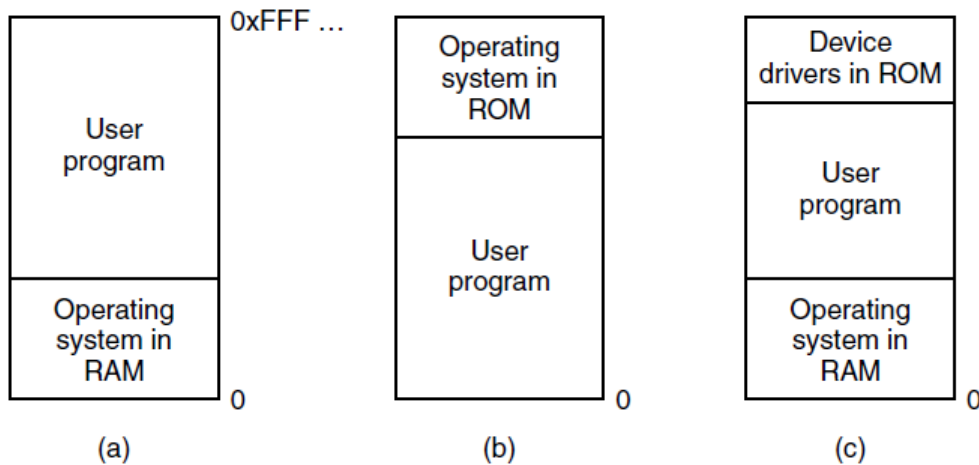
- A few megabytes of very fast, expensive, volatile cache memory
- A few gigabytes of medium-speed, medium-priced volatile main memory
- A few terabytes of slow, cheap, nonvolatile magnetic or solid state disk storage
- Some removable storages such as DVD and USB sticks

OS abstracts this hierarchy into a useful model and then manages the abstraction

# No memory abstraction

- In the past (before 1980 in PCs and before 1970 and 1960 in minicomputers and mainframe) had no memory abstraction. Every program saw the physical memory
- Under this condition, it was not possible to have two running programs in memory at the same time
  - What happen if the first program wrote a value in a physical location and then another wrote a new value to the same location???

# No memory abstraction



Three simple ways of organizing memory with an operating system and one user process.  
Other possibilities also exist

# Running multiple programs without a memory abstraction

- The early model of IBM360 divides memory into 2 KB blocks and each assigned a 4 bit protection key held in special registers inside the CPU
- A machine with 1 MB memory needed 512 of 4 bit registers for a total of 256 bytes key storage
- The IBM 360 hardware trapped any attempt by a running process to access memory with a protection code different from its key
- Only the OS could change the protection keys, user processes were prevented from interfering with other process's space

# Drawback

- Both programs refer to absolute physical memory
- Memory location in the second program (b) have to be converted
  - Static relocation was used
  - It requires extra information in all programs to indicate which would contain relocation addresses
  - Thus, it is slow and complicate

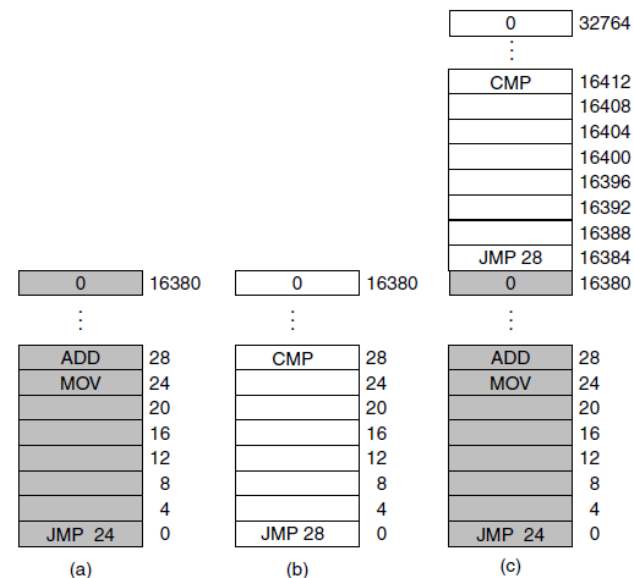


Illustration of the relocation problem. (a) A 16-KB program. (b) Another 16-KB program. (c) The two programs loaded consecutively into memory

# Memory abstraction: address space

- An address space is the set of addresses that a process can use to address memory
- Each process has its own address space
- It is independent from those belonging to other processes
- Address 28 in one program means a different physical location than address 28 in another program

# Base and limit registers

- Dynamic relocation maps each process's address space onto a different part of physical memory
- Two special hardware registers are used, "base" and "limit"
- Every time a process references memory, the CPU hardware automatically adds the base value to the address
- It checks the address is equal to or greater then the value in the "limit" register

0	32764
⋮	
CMP	16412
	16408
	16404
	16400
	16396
	16392
	16388
JMP 28	16384
0	16380

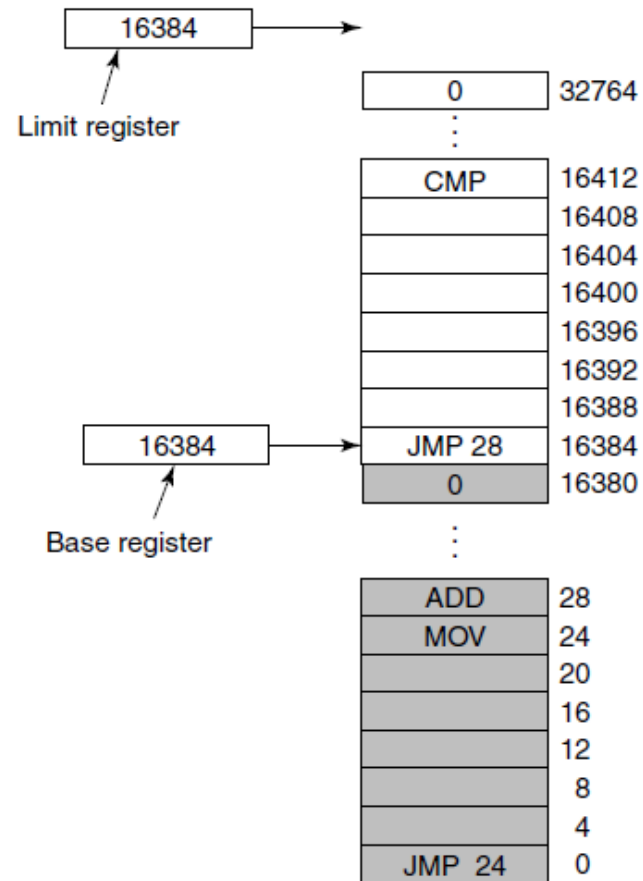
Base register contains 16384

28 is convert to 16412



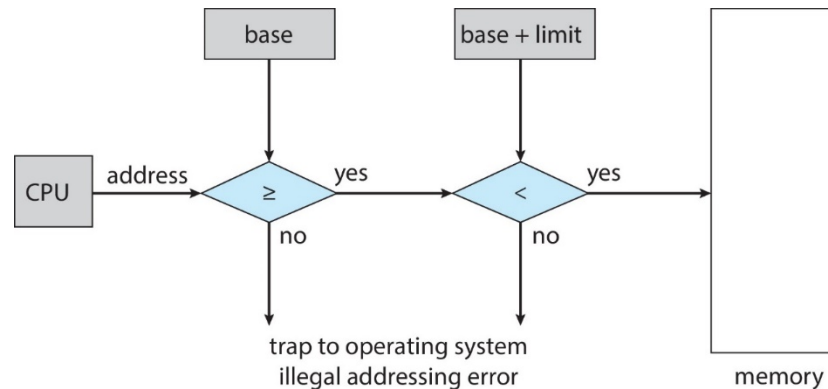
# Base and limit registers

- Base and limit registers can be used to give each process a separate address space



# Hardware address protection

- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user

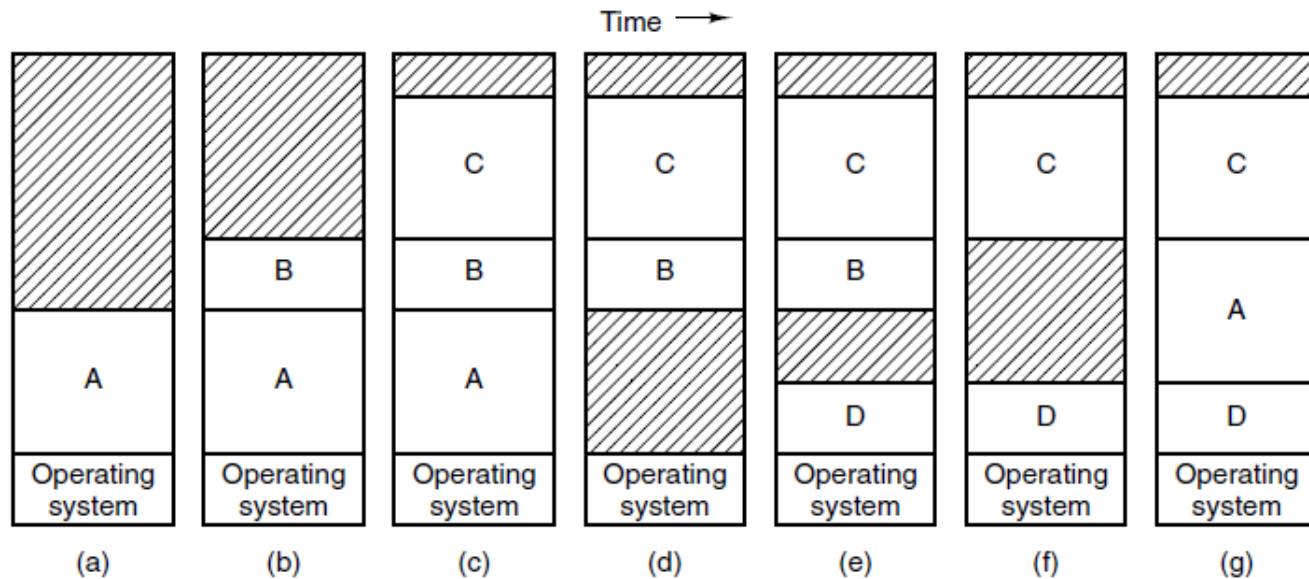


- the instructions to loading the base and limit registers are privileged by the OS

# Swapping

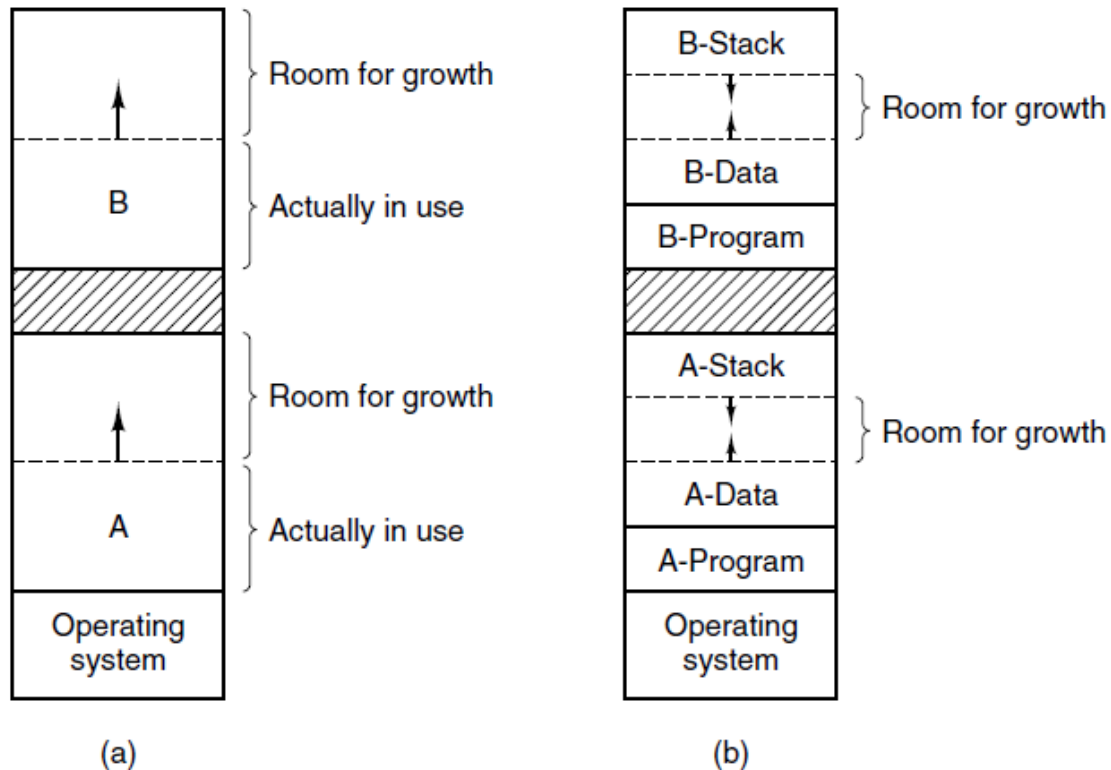
- The total amount of RAM needed by all the processes is often much more than the amount of physical memory
  - 50-100 processes may be started up as soon as the computer is booted
  - These processes may occupy 5-10 MB of memory
  - Some may require 500 MB or more
- Swapping: bringing in each process entirely in memory, running it for a while, then putting it back on the disk
- Virtual memory: allows processes to run even when they are only partially in main memory

# Swapping



Memory allocation changes as processes come into memory and leave it.  
The shaded regions are unused memory

# Memory allocation




(a) Allocating space for a growing data segment.

(b) Allocating space for a growing stack and a growing data segment

# Fixed partitions

- Physical memory is broken up into fixed partitions
  - Hardware requirements: **base register**
  - Physical address = virtual address + base register
  - Base register loaded by OS when it switches to a process
  - Size of each partition is the same and fixed
  - **How do we provide protection?**
- Advantages
  - **Easy to implement, fast context switch**
- Problems
  - **Internal fragmentation**: memory in a partition not used by a process is not available to other processes
  - **Partition size**: one size does not fit all (very large processes?)

# Fixed partitions

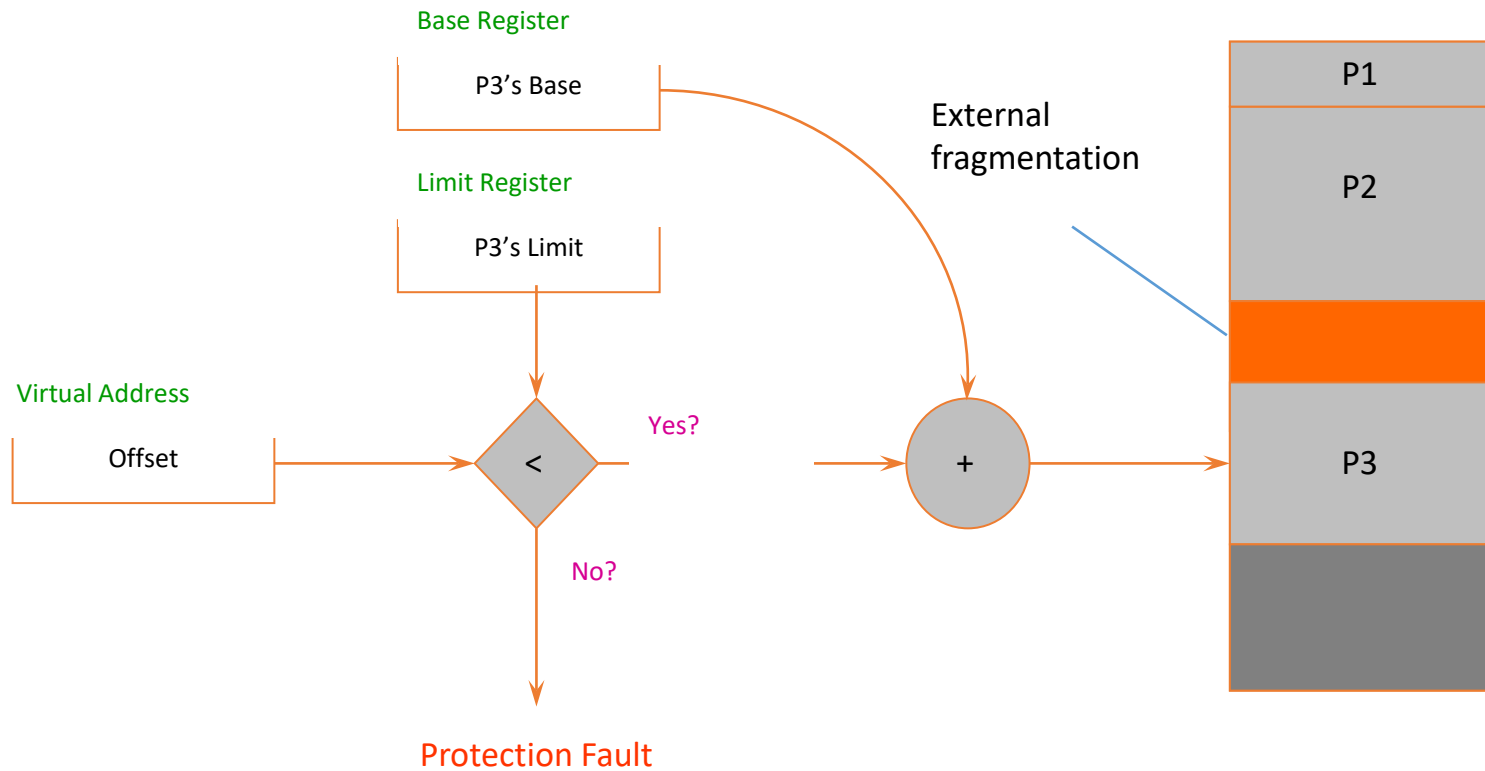


# Variable partitions

- Natural extension – physical memory is broken up into variable sized partitions
  - Hardware requirements: **base register** and **limit register**
  - Physical address = virtual address + base register
  - Why do we need the limit register? Protection
    - If (physical address > base + limit) then exception fault
- Advantages
  - **No internal fragmentation**: allocate just enough for process
- Problems
  - **External fragmentation**: job loading and unloading produces empty holes scattered throughout memory



# Variable partitions



# Variable partitions and fragmentation

Memory wasted by External Fragmentation

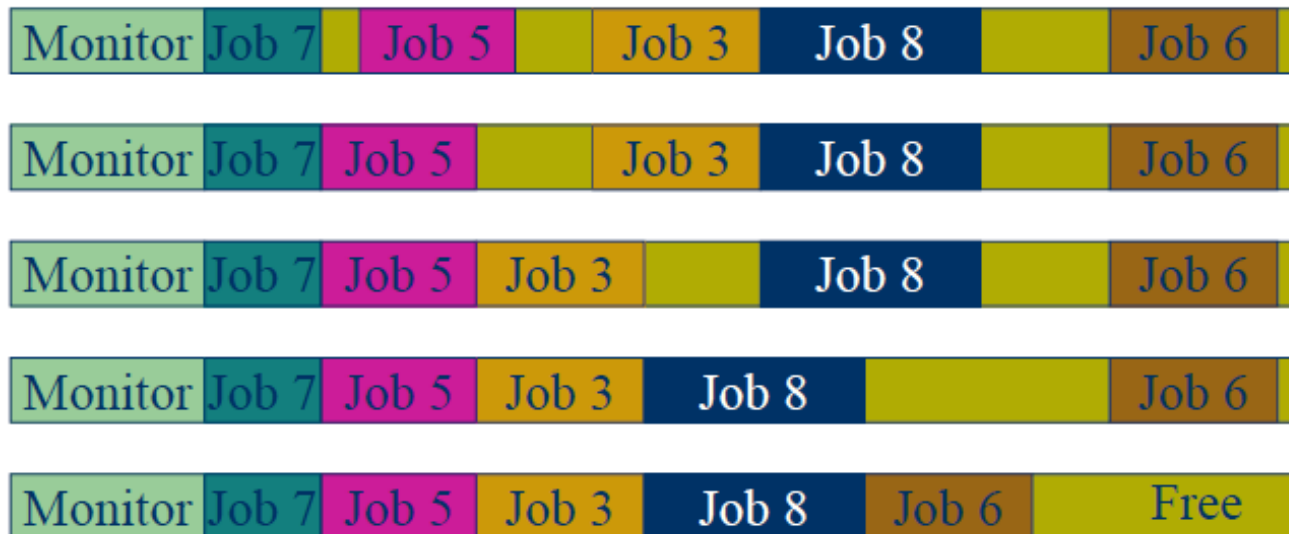


Do you know about disk de-fragmentation?

It can improve your system performance!

# Compaction

- Processes must be suspended during compaction
- Need be done only when fragmentation gets very bad



# Virtual memory

- There is a need to run programs that are too large to fit in memory
- Solution adopted in the 1960s, split programs into little pieces, called overlays
  - Kept on the disk, swapped in and out of memory
- Virtual memory : each program has its own address space, broken up into chunks called pages
- Pages are mapped onto physical memory but not all pages have to be in physical memory at the same time

# Why virtual memory?

- The abstraction that the OS will provide for managing memory is **virtual memory (VM)**
  - *Enables a program to execute with less than its complete data in physical memory*
    - A program can run on a machine with less memory than it “needs”
    - Many programs do not need all of their code and data at once (or ever) – no need to allocate memory for it
  - *Processes cannot see the memory of others*
  - OS will adjust amount of memory allocated to a process based upon its behavior
  - VM requires *hardware support* and OS management algorithms to pull it off


# Virtual addresses or logical addresses

- Virtual addresses are independent of the actual physical location of the data referenced
- OS determines location of data in physical memory
- Instructions executed by the CPU issue virtual addresses
- Virtual addresses are translated by hardware into physical addresses (with help from OS)
- The set of virtual addresses that can be used by a process comprises its virtual address space

# User/Process perspective

- Users (and processes) view memory as one contiguous address space from 0 through N
  - Virtual address space (VAS)
- In reality, pages are scattered throughout physical storage
  - Different from variable partition, where the physical memory for each process is contiguously allocated
- The mapping is **invisible** to the program
- Protection is provided because a program cannot reference memory outside of its VAS
  - The address “0x1000” maps to different physical addresses in different processes

# Memory management unit (MMU)

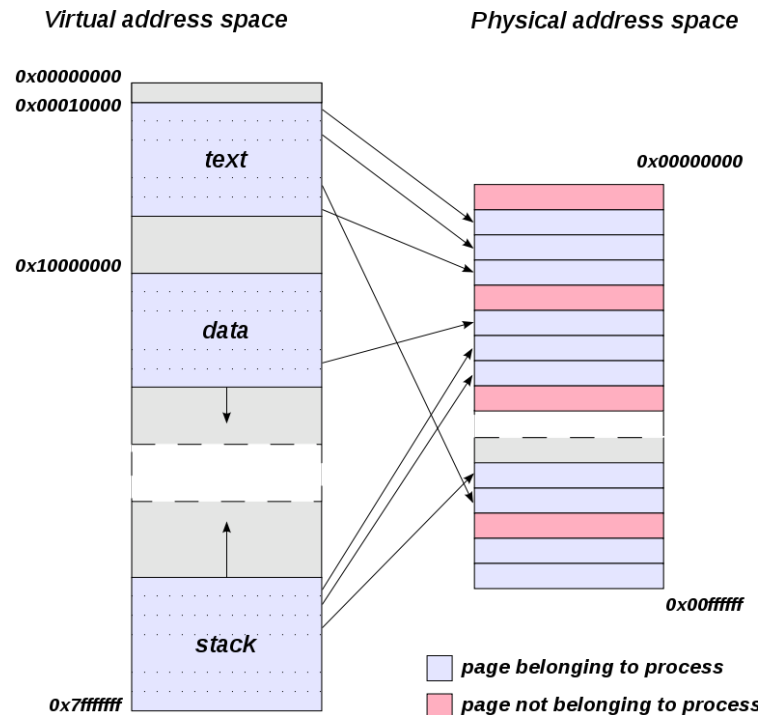


- The position and function of the MMU. Here the MMU is shown as being a part of the CPU chip because it commonly is nowadays.
- However, logically it could be a separate chip and was years ago



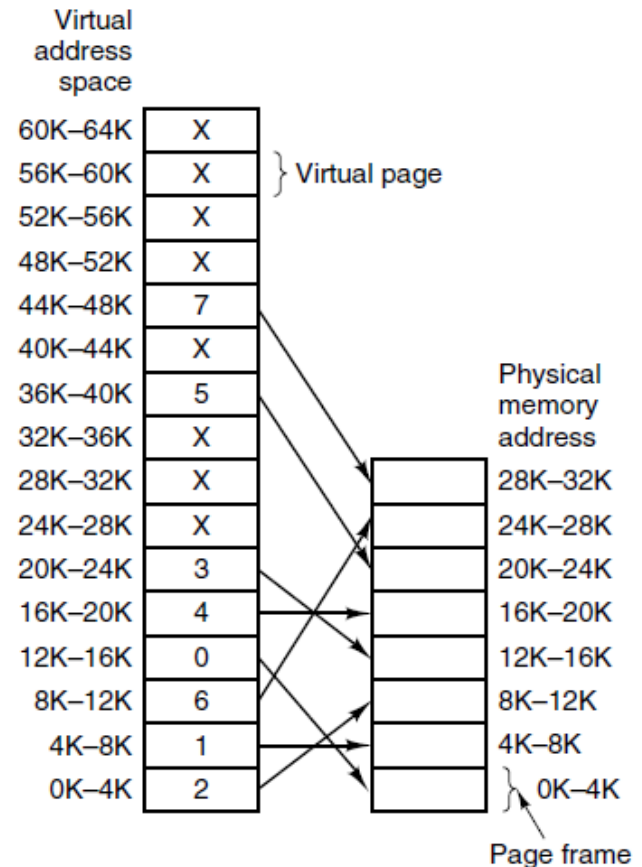
# Paging

- Paging solves the external fragmentation problem by using fixed sized units in both physical and virtual memory



# Paging

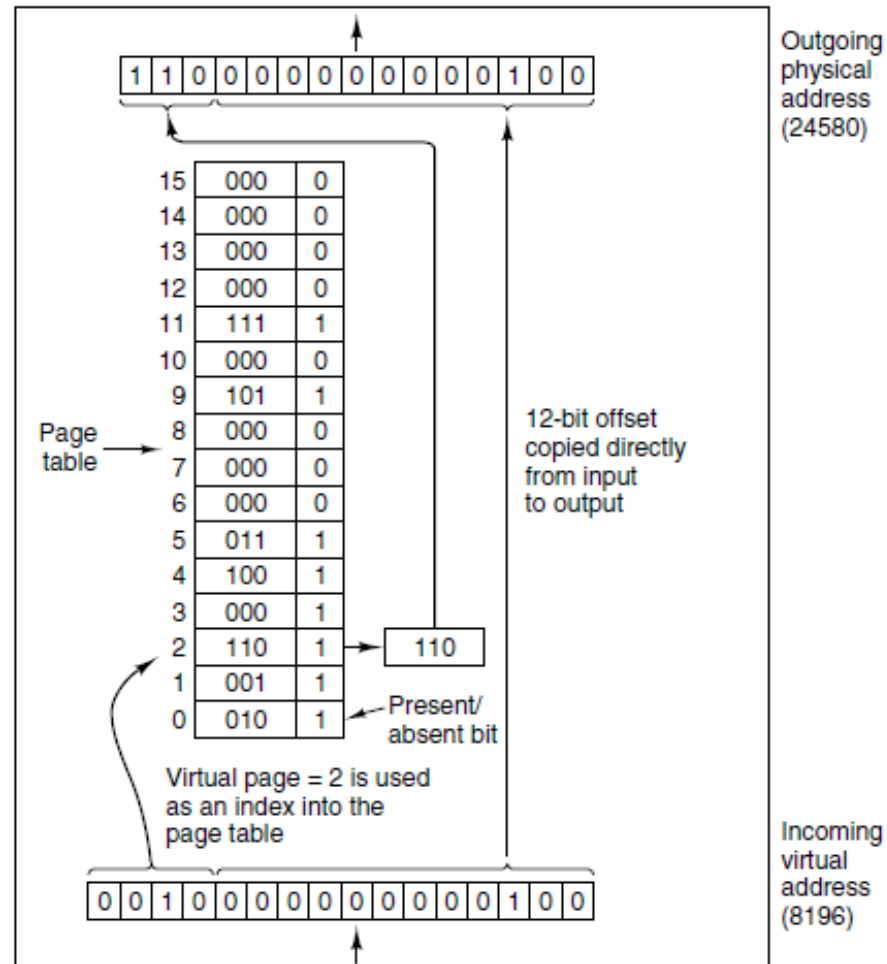
- The relation between virtual addresses and physical memory addresses is given by the page table.
- Every page begins on a multiple of 4096 and ends 4095 addresses higher,
  - so 4K–8K really means 4096–8191 and 8K to 12K means 8192–12287
  - **MOV REG, 8192**  
MMU transforms to  
**MOV REG, 24576**
- If the MMU notices that the page is unmapped (call **page fault**), OS picks a little used page and write it to the disk then fetches the missed page



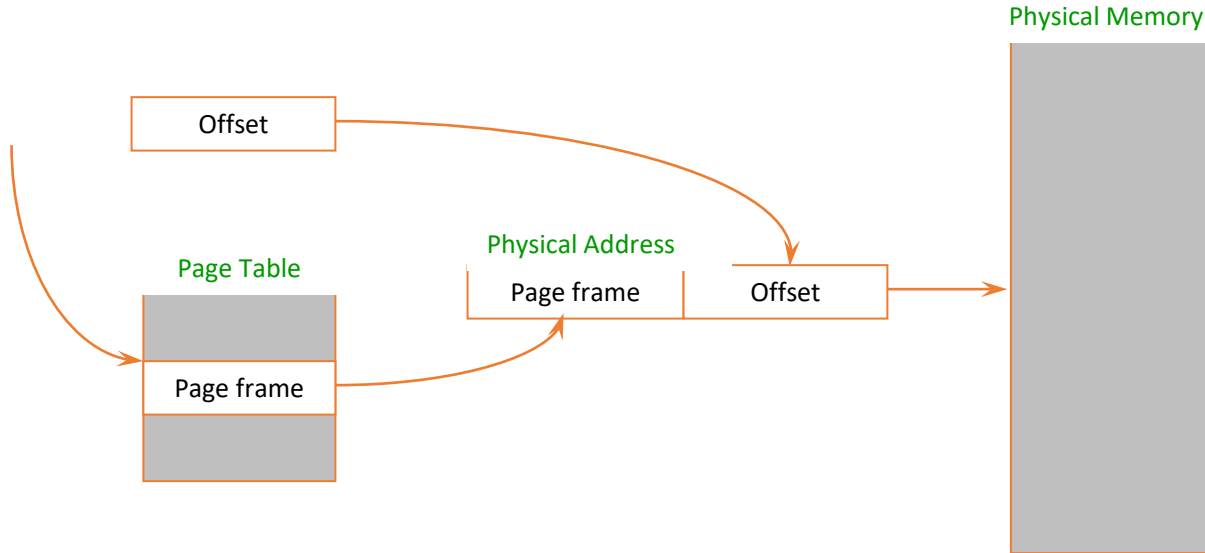
# Page translation

- Translating addresses
  - Virtual address has two parts: **virtual page number** and **offset**
  - Virtual page number (VPN) is an index into a page table
  - Page table determines page frame number (PFN)
  - Physical address is PFN::offset
- Page tables
  - Map **virtual page number** (VPN) to **page frame number** (PFN)
    - VPN is the index into the table that determines PFN
  - One page table entry (PTE) per page in virtual address space
    - Or, one PTE per VPN

# Page lookups

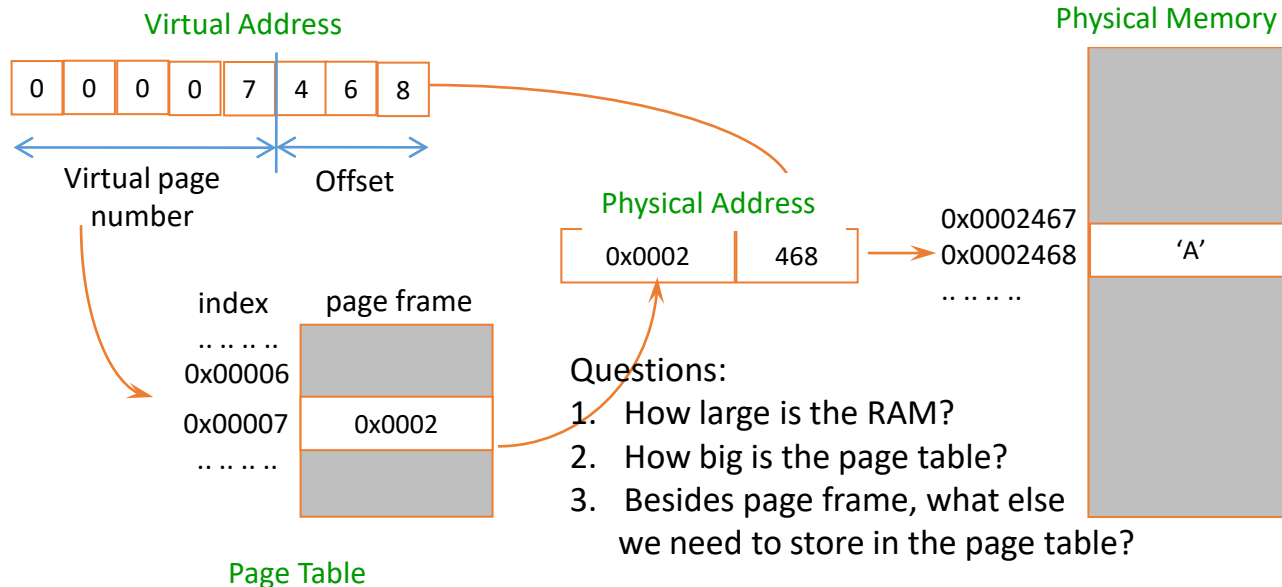


# Page lookups



# Page lookups example


Example: how do we 'load 0x00007468'?



# Page tables

- A page table is the data structure used by a virtual memory system in a computer operating system to store the mapping between virtual addresses and physical addresses
- Operating system maps the virtual address provided by the process to the physical address of the actual memory where that data is stored
- The page table is where the operating system stores its mappings of virtual addresses to physical addresses, with each mapping also known as a *page table entry*
- Each process has its own page table

# Structure of a Page Table Entry



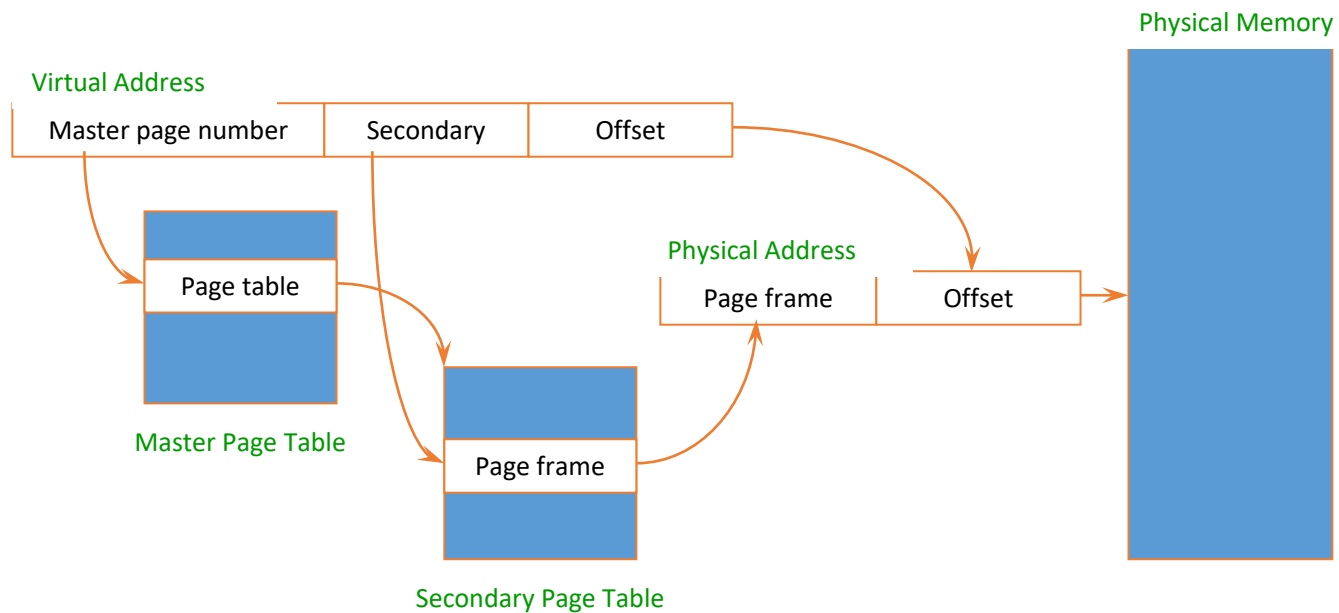
- The **Reference** bit says whether the page has been accessed
  - It is set when a read or write to the page occurs
  - The value is used to help OS choose a page to evict when a page fault occurs, pages that are not being used are better candidates to be replaced
- The **Modify** bit says whether or not the page has been written
  - It is set when a write to the page occurs
- The **Protection** bits say what operations are allowed on page
  - 0 for Read/write and 1 for read only
  - Can be 3 bits for enabling read/write/execution
- The **Present/absent** bit says whether or not the PTE can be used
  - It is checked each time the virtual address is used
  - 0 virtual page is not currently in memory, 1 the entry valid and can be used
- The **page frame number** (PFN) determines physical page



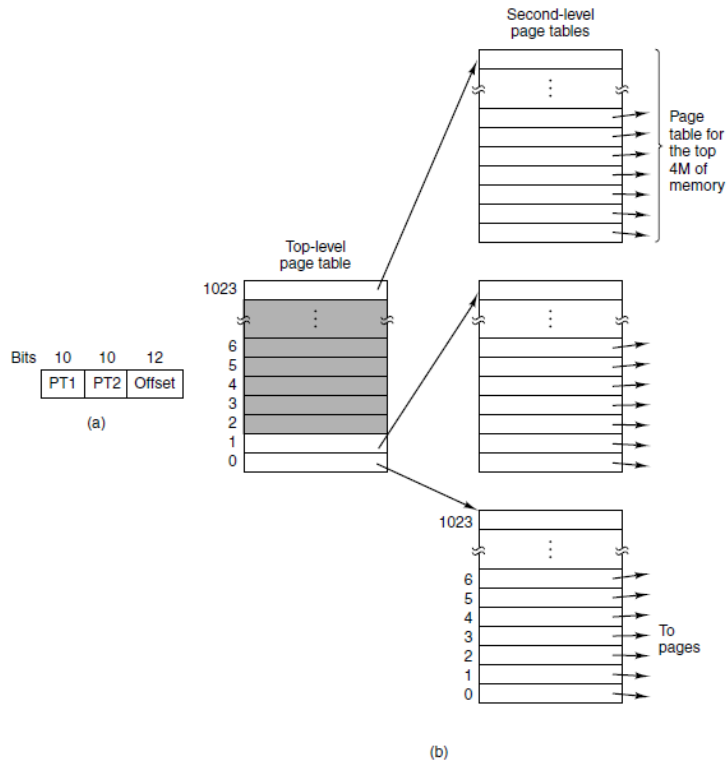
# Multi-level page table

- Single level page table size is too large
  - 4KB page, 32 bit virtual address, 1M entries per page table!
- Two-level page tables
  - Virtual addresses (VAs) have three parts:
    - Master page number, secondary page number, and offset
  - Master page table maps VAs to secondary page table
  - Secondary page table maps page number to physical page
  - Offset indicates where in physical page address is located
- Example
  - 4K pages, 4 bytes/PTE
  - How many bits in offset?  $4K = 12$  bits
  - Want master page table in one page:  $4K/4$  bytes = 1K entries
  - Hence, 1K secondary page tables. How many bits?
  - Master (1K) = 10, offset = 12, inner =  $32 - 10 - 12 = 10$  bits

# Two-Level page tables



# Two-Level page tables



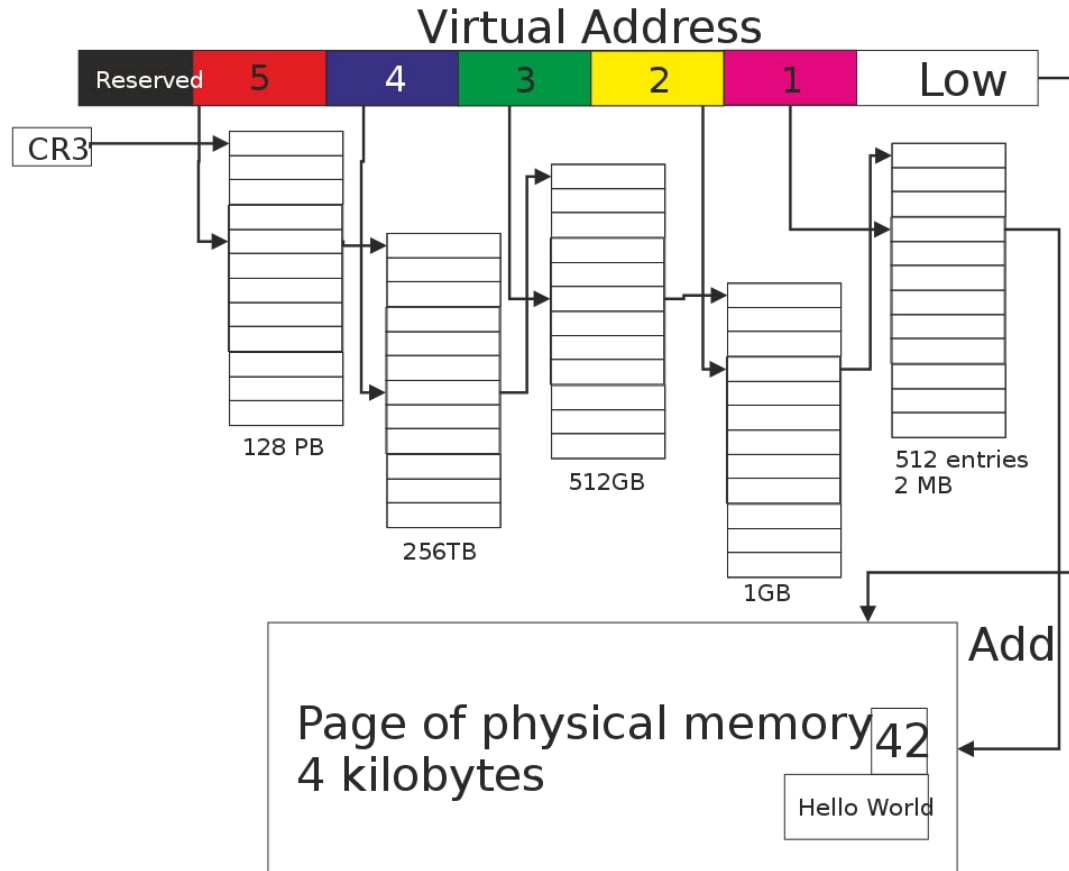
(a) A 32-bit address with two page table fields

(b) Two-level page tables

# What is the problem with 2-level page table?

- Hints:
  - Programs only know virtual addresses
  - Each virtual address must be translated
    - Each program memory access requires several actual memory accesses

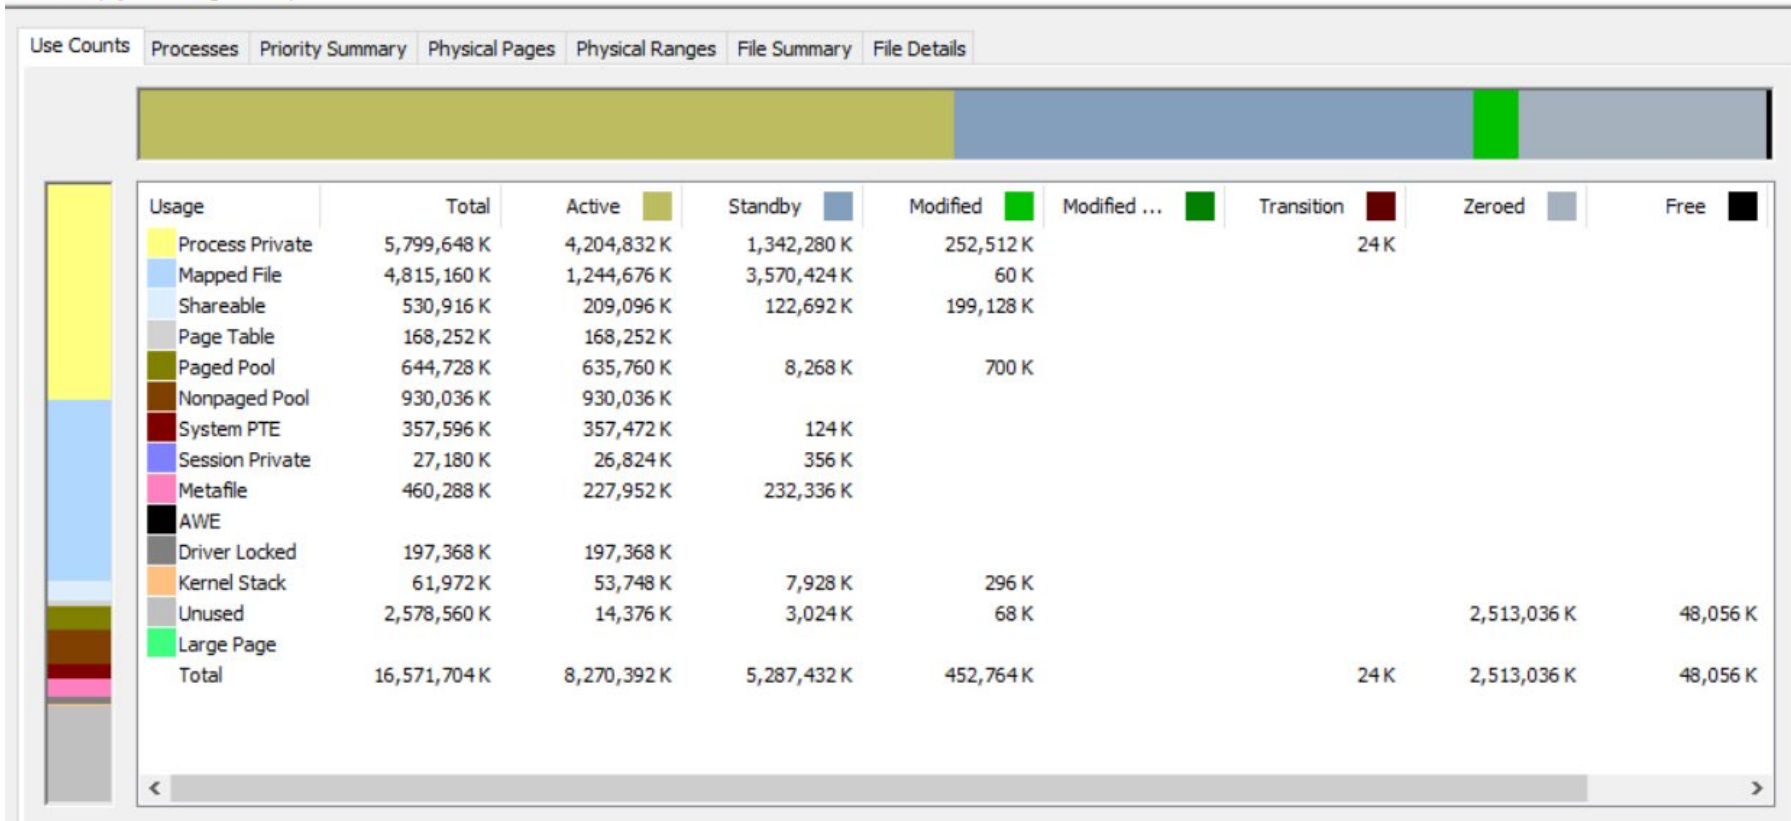
# Intel 5-level Paging in X86-64 Processor



# RamMap

RamMap - Sysinternals: [www.sysinternals.com](http://www.sysinternals.com)

File Empty Settings Help



# Paging advantages

- Easy to allocate memory
  - Memory comes from a free list of fixed size chunks
  - Allocating a page is just removing it from the list
  - External fragmentation not a problem
- Easy to **swap** out chunks of a program
  - All chunks are the same size
  - Use valid bit to detect references to swapped pages
  - Pages are a convenient multiple of the disk block size

# Paging limitations

- Can still have internal fragmentation
  - Process may not use memory in multiples of a page
- Memory reference overhead
  - 2 references per address lookup (page table, then memory)
    - Even more for two-level page tables!
  - Solution – use a hardware cache of lookups
- Memory required to hold page table can be significant
  - Need one PTE per page
  - 32 bit address space w/ 4KB pages =  $2^{20}$  PTEs
  - 4 bytes/PTE = 4MB/page table
  - 25 processes = 100MB just for page tables!
    - Remember: each process has its own page table!
  - Solution – 2-level page tables



# What if a process requires more memory than physical memory?

- Swapping
  - Move one/several/all pages of a process to **disk**
    - Free up physical memory
    - “Page” is the unit of swapping
  - The freed physical memory can be mapped to other pages
  - Processes that use large memory can be swapped out (and later back in)

# A variation of paging: Segmentation

- Segmentation is a technique that partitions memory into logically related data units
  - Module, procedure, stack, data, file, etc.
  - Virtual addresses become <segment #, offset>
  - Units of memory from user's perspective
- Natural extension of variable-sized partitions
  - Variable-sized partitions = 1 segment/process
  - Segmentation = many segments/process
- Hardware support
  - Multiple base/limit pairs, one per segment (segment table)
  - Segments named by #, used to index into table

# Segment table

- Extensions
  - Can have one segment table per process
    - Segment #s are then process-relative
  - Can easily share memory
    - Put same translation into base/limit pair
    - Can share with different protections (same base/limit, diff prot)
- Problems
  - Large segment tables
    - Keep in main memory, use hardware cache for speed
  - Large segments
    - Internal fragmentation, paging to/from disk is expensive

# Segmentation and Paging

- Can combine segmentation and paging
  - The x86 supports segments and paging
- Use segments to manage logically related units
  - Module, procedure, stack, file, data, etc.
  - Segments vary in size, but usually large (multiple pages)
- Use pages to partition segments into fixed size chunks
  - Makes segments easier to manage within physical memory
    - Segments become “pageable” – rather than moving segments into and out of memory, just move page portions of segment
  - Need to allocate page table entries only for those pieces of the segments that have themselves been allocated
- Tends to be complex...

# Summary

- Virtual memory
  - Processes use virtual addresses
  - OS + hardware translates virtual address into physical addresses
- Various techniques
  - Fixed partitions – easy to use, but internal fragmentation
  - Variable partitions – more efficient, but external fragmentation
  - Paging – use small, fixed size chunks, efficient for OS
  - Segmentation – manage in chunks from user's perspective
  - Combine paging and segmentation to get benefits of both

# Sources and References

- Andrew S. Tanenbaum, Herbert Bos, Modern Operating 4<sup>th</sup> edition, Pearson, 2015
- Presentations by Ding Yuan, ECE Dept., University of Toronto
- Operating System Concepts 10<sup>th</sup> book official slides by Abraham Silberschatz, Greg Gagne, Peter B. Galvin